

Data and Programming Analytics

Yelp Review Sentiment Analysis

Group 18

Yaohui Wu, Jai Agrawal, Jennie Ryu, Xiaofan Zhu, Sahithi Muddana



Table of Content

1 Introduction

- Market
- Project Task
- Data Introduction
- Pre-processing

2 Analysis

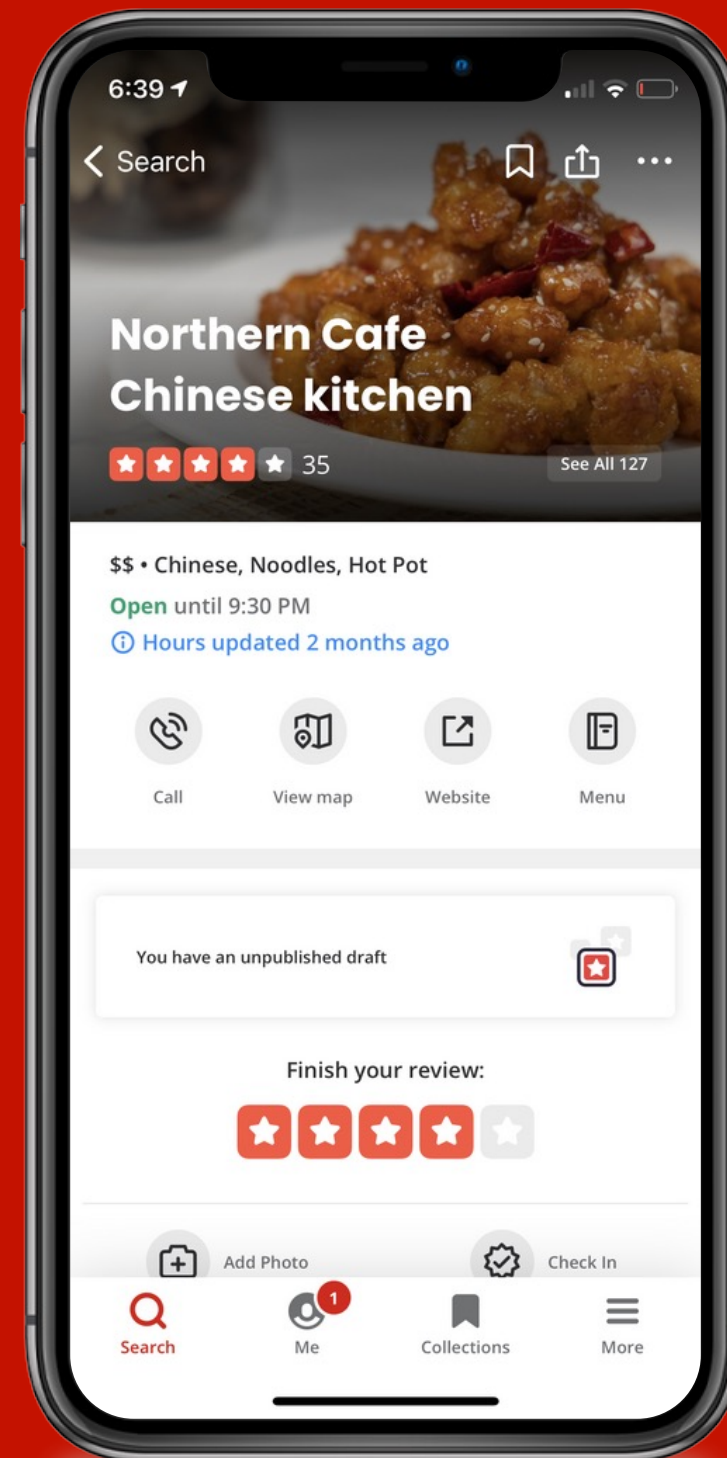
- Sentiment Analysis
- Topic Modeling

3 Conclusion

- Conclusion
- Q&A

2021 Q2 revenue is more than

\$257M



Yelp rating can increase
a business owner's sales by

5-9%

PROJECT TASK



Topic modeling (Positive)

Topic modeling to extract the reasons why people give positive comments to business.



Topic modeling (Negative)

Topic modeling to extract the reasons why people give negative comments to business.



Topic modeling (Neutral)

Sentiment Analysis and topic modeling extract actual negative and positive comments from neutral comments

DATA INTRODUCTION

business.json

- business_id
- stars
- city
- state
- review_count.

Review.json

- review_id
- business_id
- stars
- useful
- etc
- Dataset size: 8.6 millions

```
with open('drive/MyDrive/Colab Notebooks/yelp_dataset/yelp_academic_dataset_business.json') as f:
    for i in f:
        line = eval(i.replace('null', '" null"', 100))
        if count == 0:
            if 'business_id' not in business_info.keys():
                business_info['business_id'] = [line.get('business_id')]
            else:
                business_info['business_id'].append(line.get('business_id'))
            if 'name' not in business_info.keys():
                business_info['name'] = [line.get('name')]
            else:
                business_info['name'].append(line.get('name'))
            if 'address' not in business_info.keys():
                business_info['address'] = [line.get('address')]
            else:
                business_info['address'].append(line.get('address'))
            if 'city' not in business_info.keys():
                business_info['city'] = [line.get('city')]
            else:
                business_info['city'].append(line.get('city'))
            if 'state' not in business_info.keys():
                business_info['state'] = [line.get('state')]
            else:
                business_info['state'].append(line.get('state'))
            if 'postal_code' not in business_info.keys():
                business_info['postal_code'] = [line.get('postal_code')]
            else:
                business_info['postal_code'].append(line.get('postal_code'))
            if 'stars' not in business_info.keys():
                business_info['stars'] = [line.get('stars')]
            else:
                business_info['stars'].append(line.get('stars'))
            if 'review_count' not in business_info.keys():
                business_info['review_count'] = [line.get('review_count')]
            else:
                business_info['review_count'].append(line.get('review_count'))
```


RESERVOIR SAMPLING ALGORITHM

```
1 #reservoir sampling algorithm code
2 def reservoir_sampling(sampled_num, total_num):
3     pool = []
4     for i in range(0, total_num):#i is from 0 to 100000-1
5         if i < sampled_num:
6             pool.append(i)
7         else:
8             r = random.randint(0, i)
9             if r < sampled_num:
10                 pool[r] = i
11     return pool
```

1

Reservoir Sampling Method

2

Randomly pick 100,000 value from 0 to 8,635,403

3

Use as Index

RESERVOIR SAMPLING ALGORITHM

```
1 #pool is a 100000 length list contains value from 0 to 8635402
2 import datetime
3 start = datetime.datetime.now()
4 if 8665403 in pool:
5     print('a')
6 end = datetime.datetime.now()
7 print(end-start)
```

0:00:00.011099

```
1 pool_hash = set(pool)
2 #pool_hash is a 100000 length list contains value from 0 to 8635402
3 import datetime
4 start = datetime.datetime.now()
5 if 8665403 in pool_hash:
6     print('a')
7 end = datetime.datetime.now()
8 print(end-start)
```

0:00:00.000104

```
1 print(0.011099/0.000104)
```

106.72115384615385

20 Hours



Hashing

20 Minutes

DATA VISUALIZAION

```
1 df.info()
```


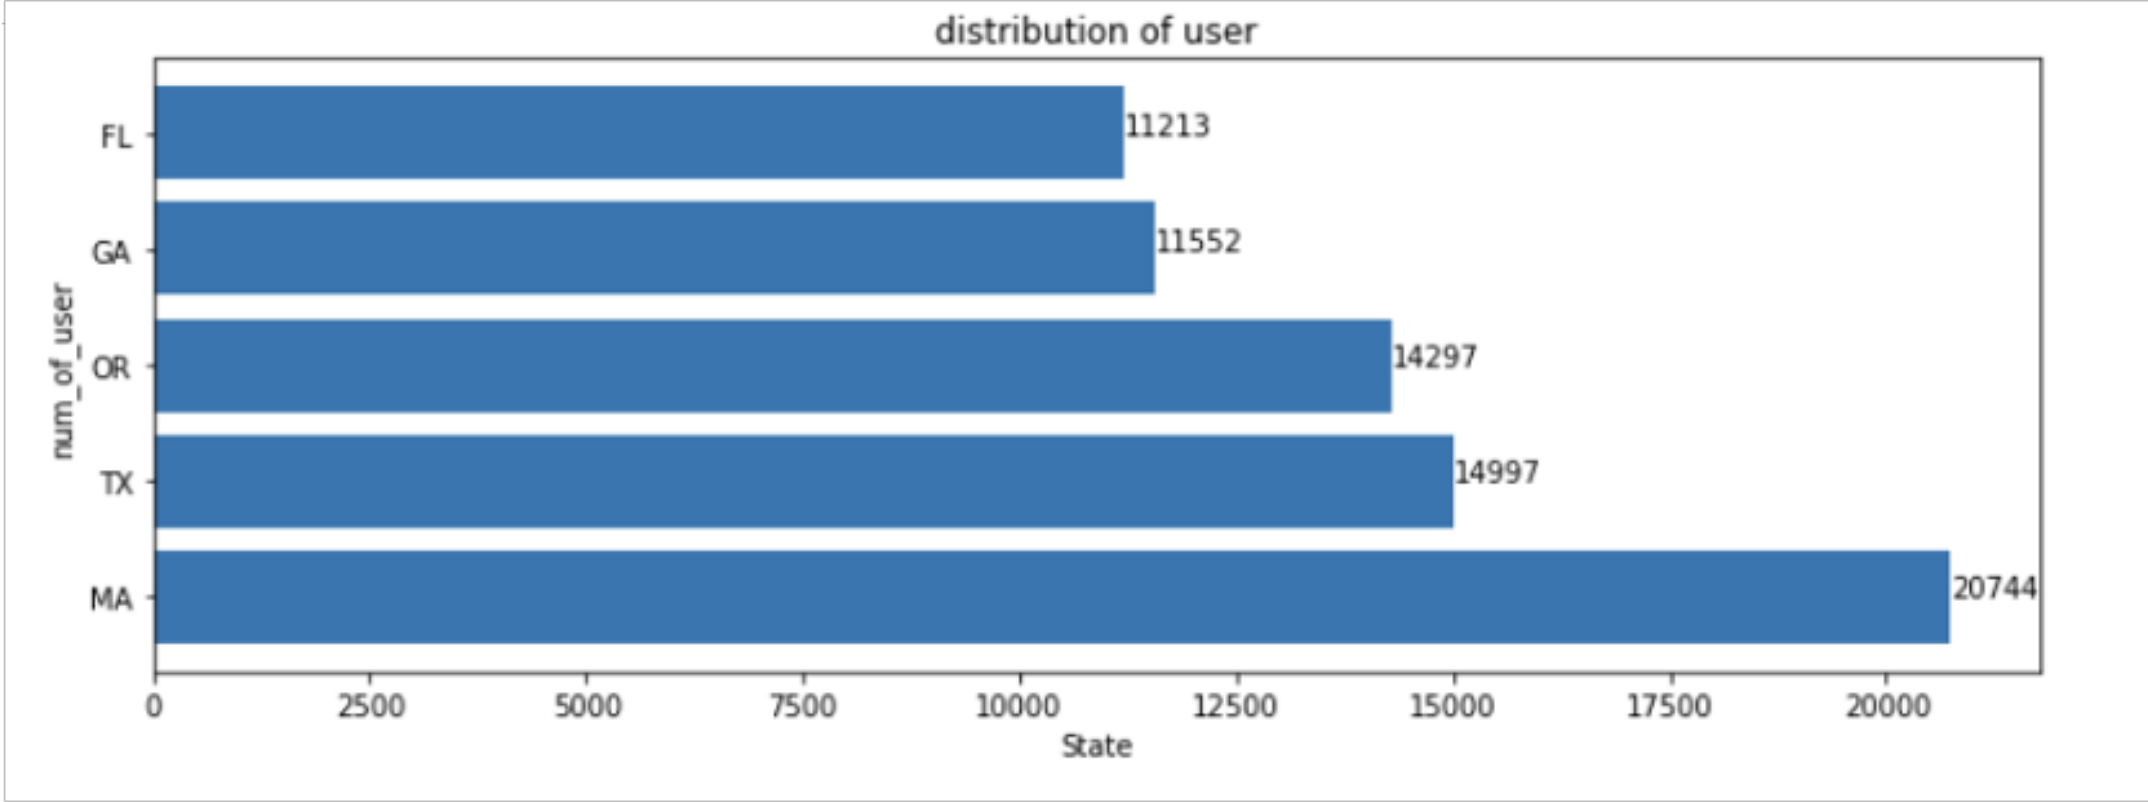

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100000 entries, 0 to 99999
Data columns (total 16 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   review_id       100000 non-null  object 
 1   user_id         100000 non-null  object 
 2   business_id     100000 non-null  object 
 3   stars           100000 non-null  float64 
 4   useful          100000 non-null  int64   
 5   funny           100000 non-null  int64   
 6   cool            100000 non-null  int64   
 7   text            100000 non-null  object 
 8   date            100000 non-null  object 
 9   name            100000 non-null  object 
10  address         100000 non-null  object 
11  city            100000 non-null  object 
12  state           100000 non-null  object 
13  postal_code     100000 non-null  object 
14  avg_stars       100000 non-null  float64 
15  review_count    100000 non-null  int64   
dtypes: float64(2), int64(4), object(10)
memory usage: 13.0+ MB
```

100,000 Rows

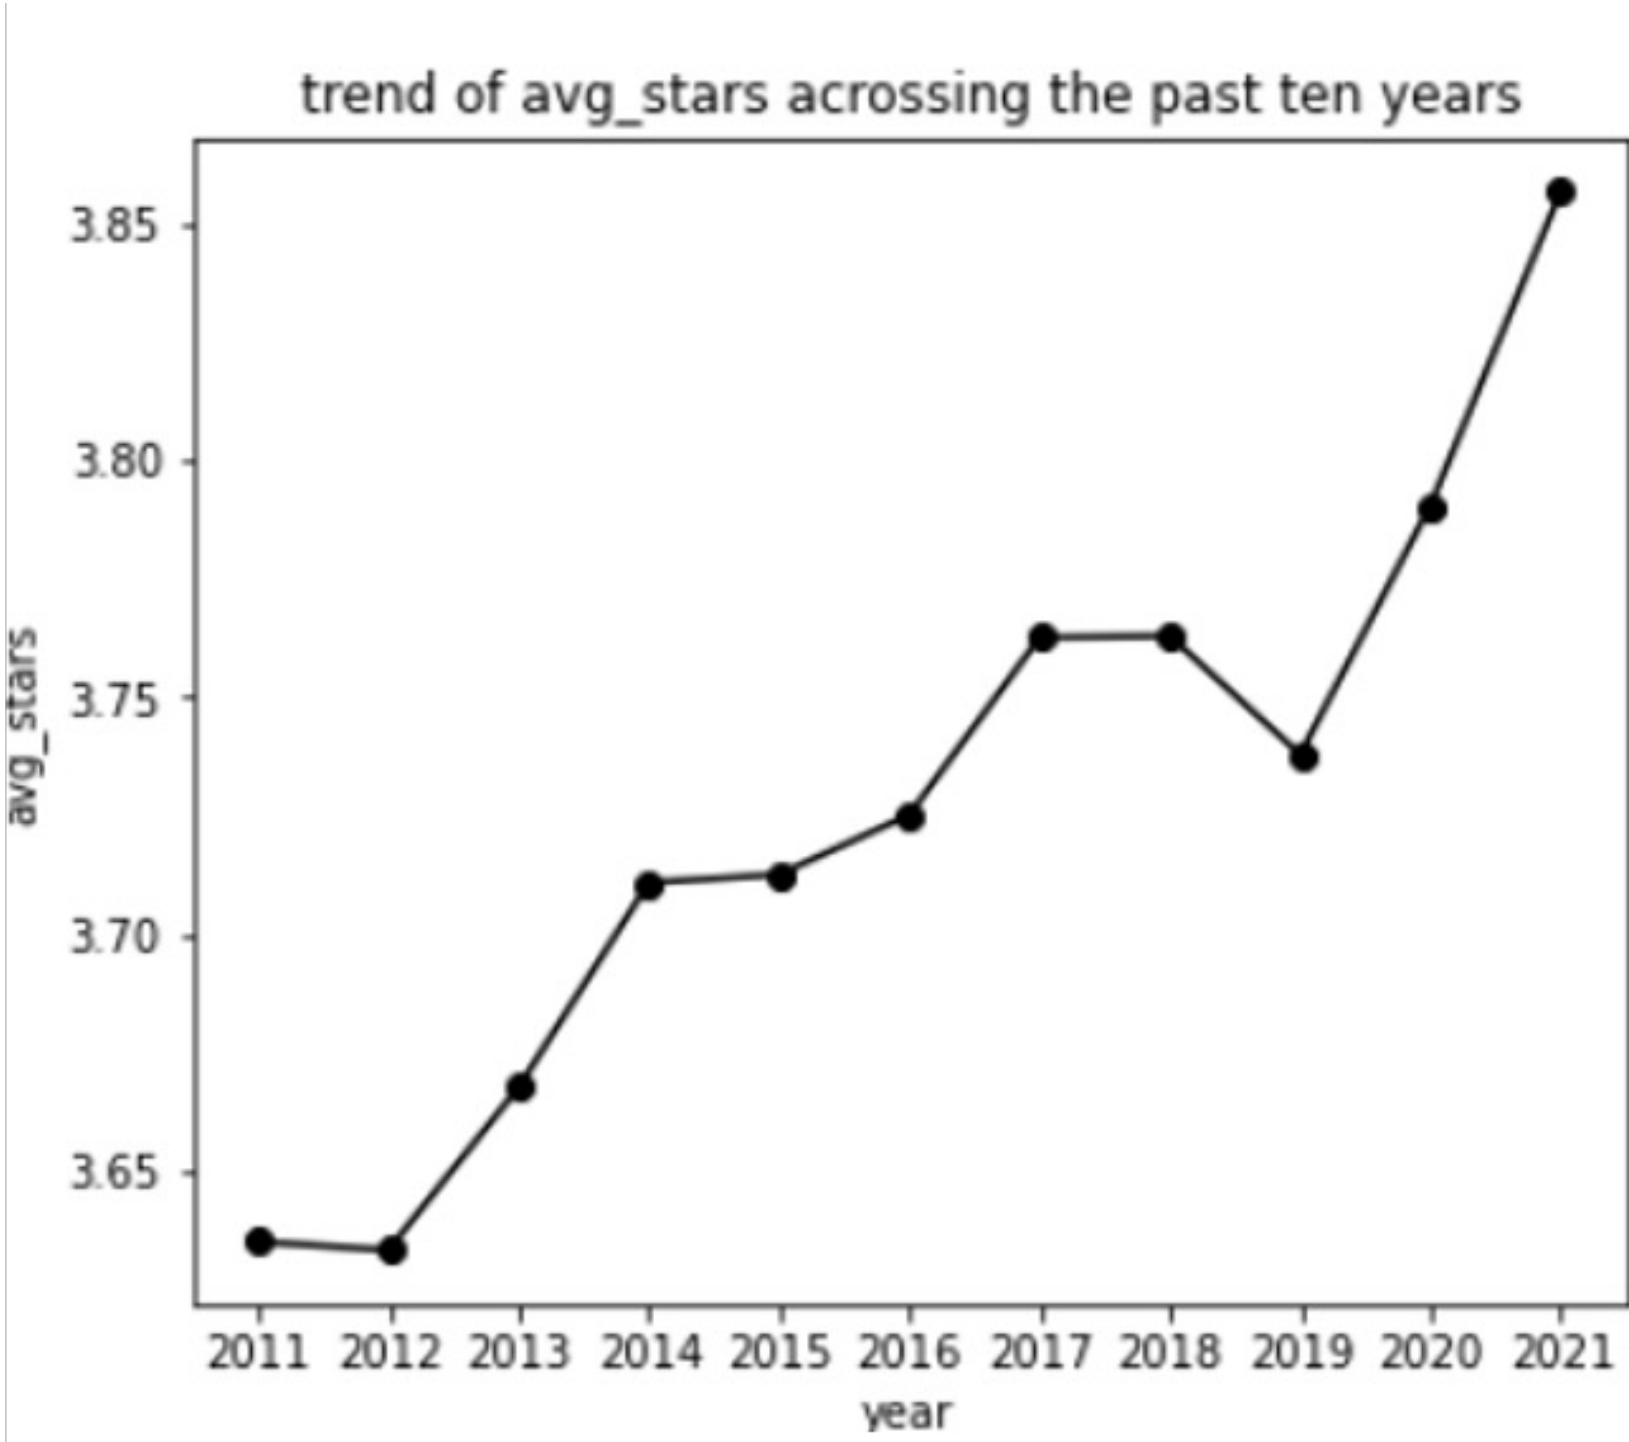
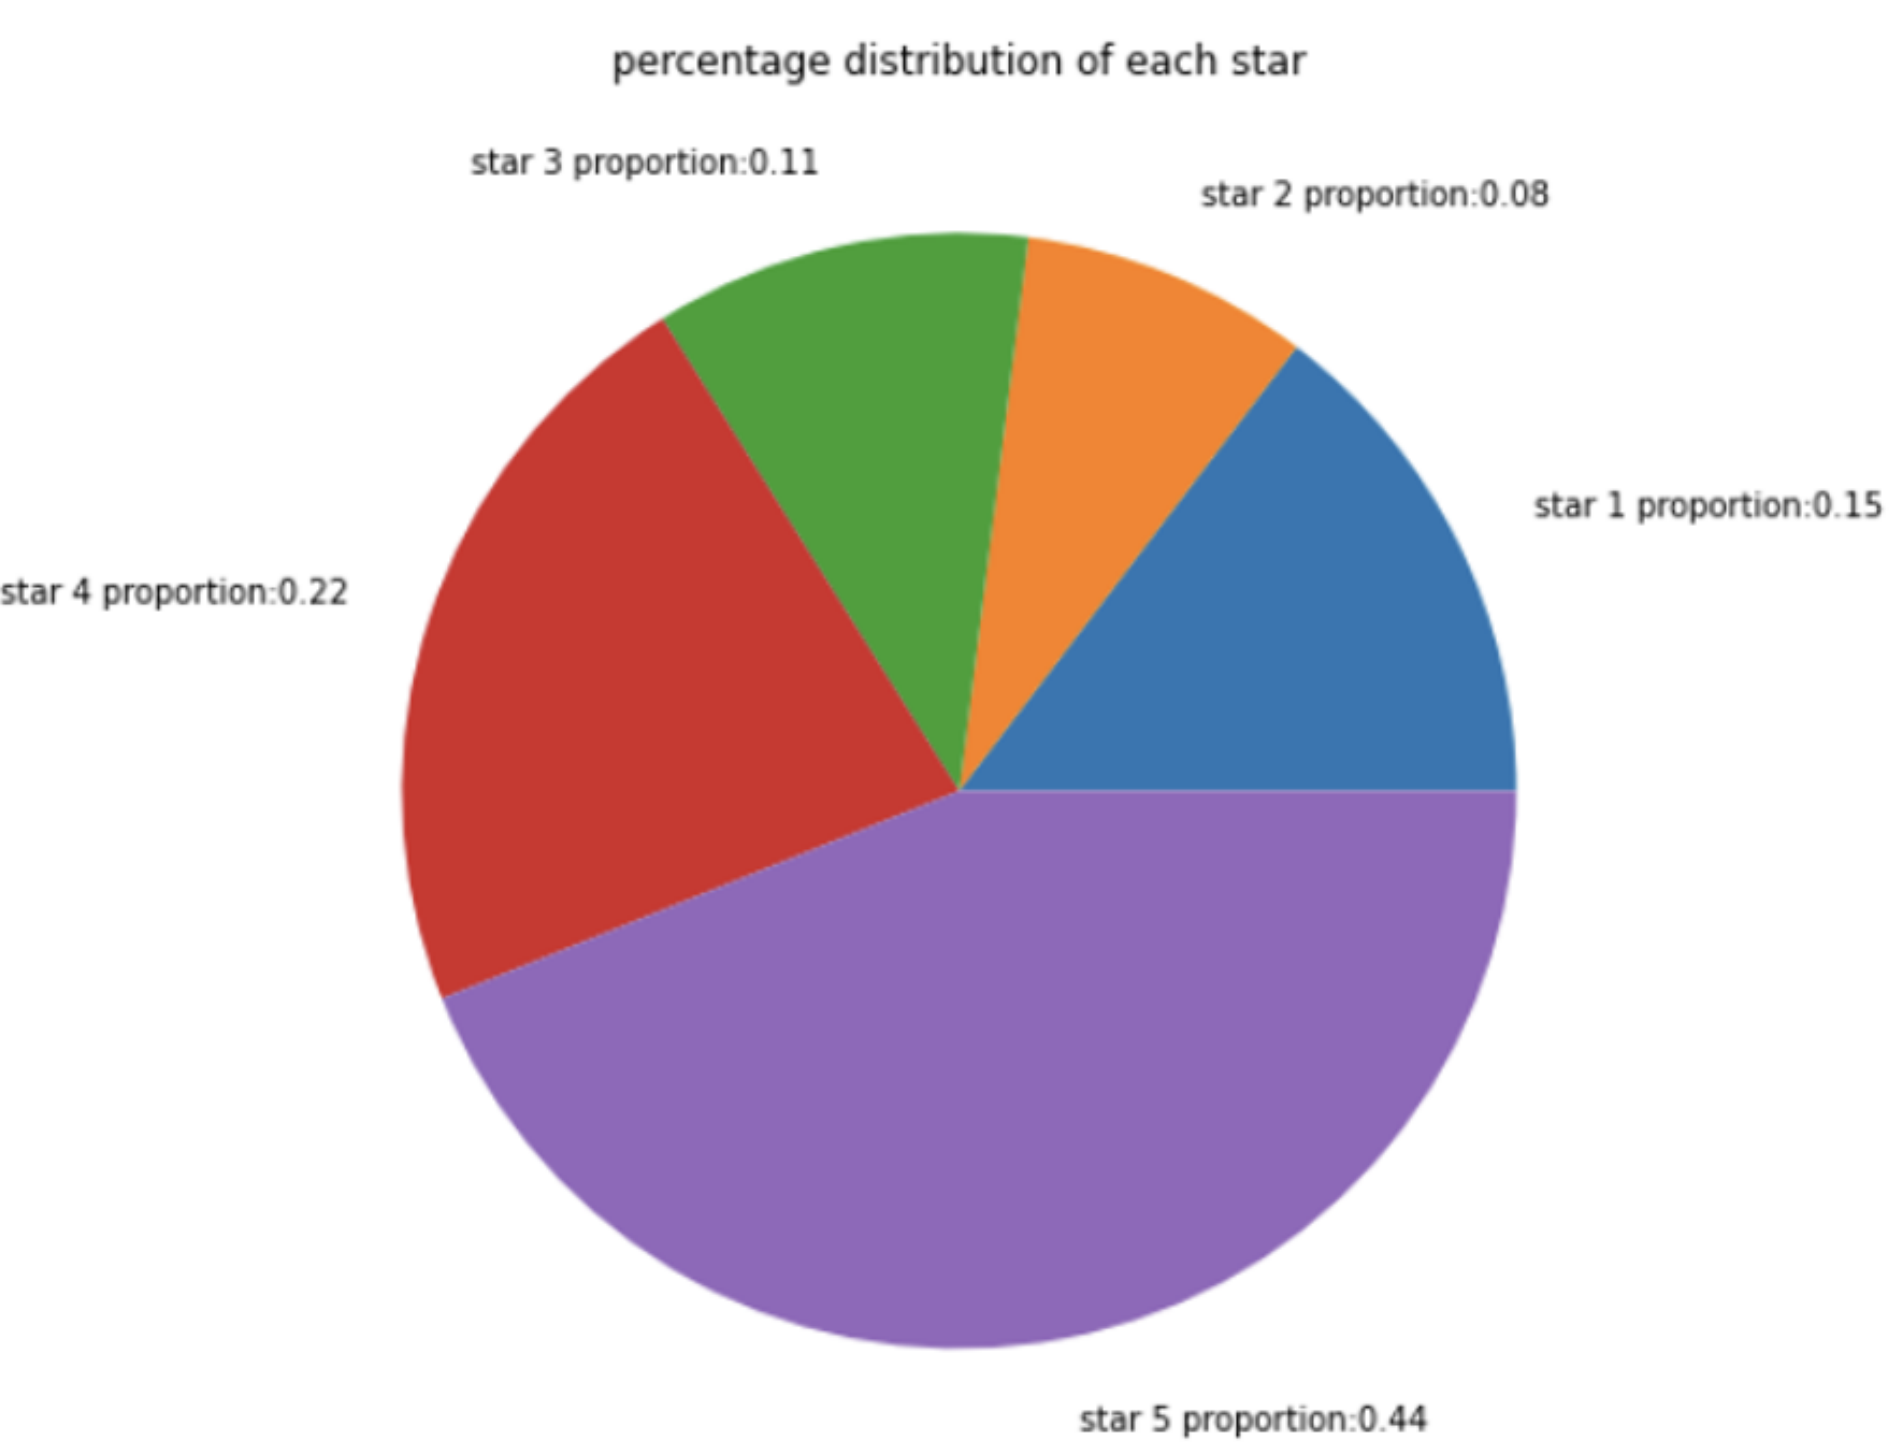
16 Columns

0 N/A Values

GA(Georgia),FL(Florida),TX(Texas),OR(Oregon),MA(Massachusetts) are the top five states



positive(4 or 5 stars) comments takes 66%, negative(1 or 2 stars) take 23%, neutral(3 stars) takes 11%



DATA CLEANING

Remove all of the reviews in foreign languages

```
807      デイビッドのローカルなフローズンヨーグルト屋さん。自前のキャラも作って、お店も広くて清潔。...
1075     Solid modern Shanghainese food. Quality ingred...
2218     店員のおねーちゃんがタイ人かな？とっても笑顔が素敵でした。お水もなくなったらすぐ入れに来てく...
2627     一月十四日我們四人叫了菜，其中一道油豆腐已經壞了，有告知服務人員，他說另外給我們一份油豆腐，...
3687     20180504 after my interview, I went downtown V...

...
92533     没想到吃火锅会食物中毒！我和先生前天晚上10月11号晚5点后到达火锅店，以前也来过这家两次...
93090     Restaurant was rather packed and required some...
94282     小さな屋台である。希望を聞いて巻いてくれるが、スパムと野菜を巻いてもらって美味しかった。ソウ...
97901     We really like their fish bbq 烤鱼 and also real...
99499     Among all, Kung Fu (功夫茶 is probably the sounde...
Name: text, Length: 64, dtype: object
```

DATA PRE-PROCESSING

- **Make all text "lower case"**
 - **Remove "stopwords" from text**
 - **"Amplify" the voice of each user's comment**
-

SENTIMENT ANALYSIS



NEGATIVE

1~2 stars




NEUTRAL



POSITIVE

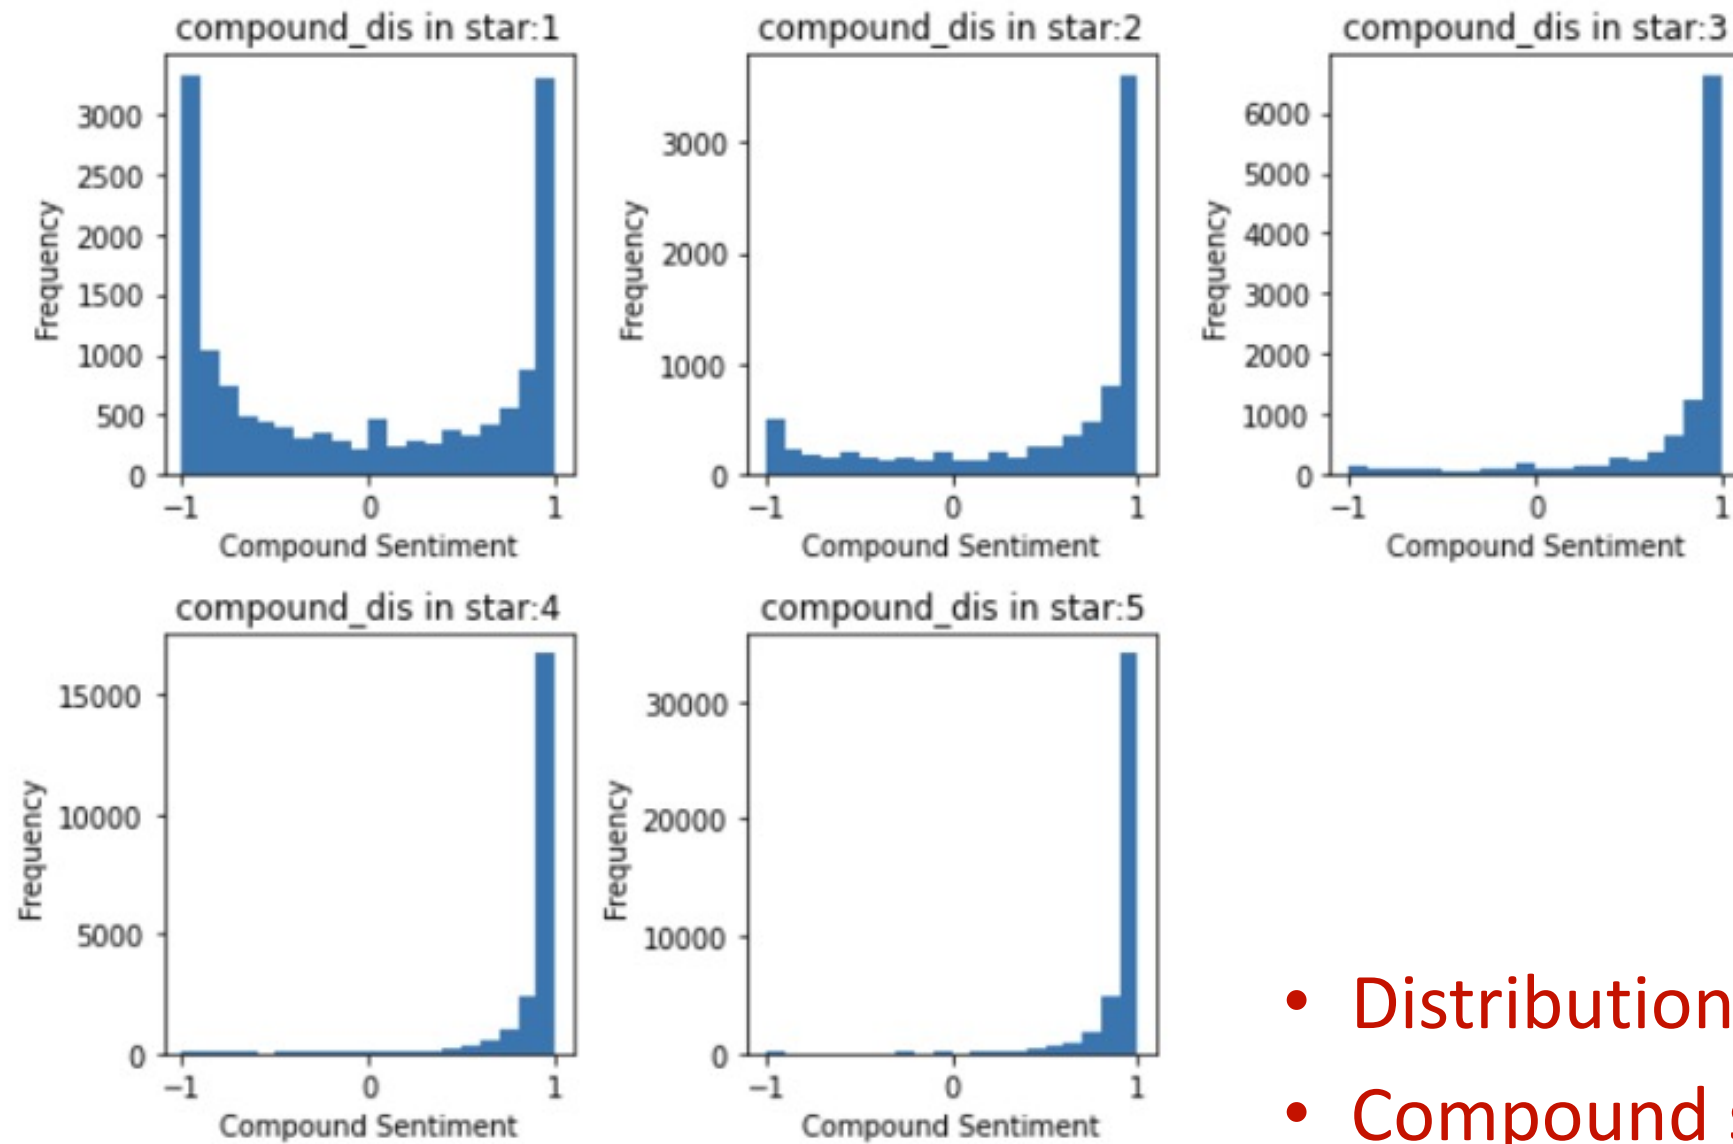
4~5 stars

SENTIMENT ANALYSIS



- Average stars of the dataset is around 3.5
- Compound score is skewed to positive 1

SENTIMENT ANALYSIS



- Distribution of compound score change when star increases
 - Compound score moved to positive 1
-

TOPIC MODELING

Model Building

Create our own stopwords package

- Stopword(English) From spacy.lang.en.stop_words: 326
- Stopword(English) From NLTK: 179
- Manually created stopwords: 20
- Our own stopwords: 525

```
my_stopwords = ['good', 'great', 'bad', 'recommend', 'like', 'pizza', 'chicken', 'horrible', 'better', 'worst', 'worse',  
                'really', 'best', 'amazing', 'excellent', 'really', 'nice', 'love', 'highly', 'pretty']  
  
# stopwords from spacy  
from spacy.lang.en.stop_words import STOP_WORDS as en_stop  
  
# stopwords from nltk  
stopword_forTfidfVectorizer = list(stopwords.words('english')) + list(en_stop) + my_stopwords
```

TOPIC MODELING

```
vectorizer = TfidfVectorizer(stop_words=stopword_forTfidfVectorizer,  
                             #stop_words='english',  
                             max_features= 500 # keep top 1000 terms  
                             ,max_df = 0.35  
                             ,min_df = 0.05  
                             )
```

(A, B)	C
(0, 125)	0.34835111471529684
(0, 112)	0.3580215090225704
(0, 32)	0.3984352082029299
(0, 34)	0.43893435333429925
(0, 24)	0.4091533930501054
(0, 78)	0.4813020041893511
(1, 111)	0.20651731665525333
(1, 110)	0.3498094344052366
(1, 9)	0.2609792801954902
(1, 3)	0.33487668065116805
(1, 30)	0.3608773291313086
(1, 88)	0.31003361638379445
(1, 107)	0.3375006685332691
(1, 47)	0.3316450335079065

TfidfVectorizer

- A : Index for Comment
- B : Index for Keyword
- C : Level of Importance

TOPIC MODELING

```
svd_model = TruncatedSVD(n_components=topicsNum, algorithm='randomized', n_iter=100, random_state=122)
```

```
components_ output shape (10, 133)
[[ 0.05375653  0.05193305  0.04573669 ...  0.0534104  0.05281103
   0.05644966]
 [-0.01670479 -0.00139009 -0.00413807 ...  0.01598738 -0.01965294
  -0.04217895]
 [ 0.02186757  0.03559468 -0.01549099 ...  0.02746867 -0.0054664
   0.00959735]
 ...
 [ 0.01222876 -0.01604516 -0.04206287 ...  0.03657012  0.05360057
   0.04330051]
 [ 0.00607374  0.00245995  0.01883561 ...  0.03427292 -0.03636514
  -0.00046736]
 [ 0.01498355  0.03766232 -0.00737106 ...  0.04572849  0.01750786
   0.02769219]]
```

TruncatedSVD

- Row : Topic
- Value : Importance of the keyword

TOPIC MODELING


```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import TruncatedSVD
def getTopics(df, topicsNum):
    re = ''
    vectorizer = TfidfVectorizer(stop_words=stopword_forTfidfVectorizer,
                                #stop_words='english',
                                max_features= 500 # keep top 1000 terms
                                ,max_df = 0.35
                                ,min_df = 0.05
                                )

    X = vectorizer.fit_transform(df['amplified_text'])
    svd_model = TruncatedSVD(n_components=topicsNum, algorithm='randomized', n_iter=100, random_state=122)
    svd_model.fit(X)


    terms = vectorizer.get_feature_names()
    print(len(terms))
    print(svd_model.components_.shape)
    print(svd_model.feature_names_in_.shape)

    for i, comp in enumerate(svd_model.components_):
        terms_comp = zip(terms, comp)
        sorted_terms = sorted(terms_comp, key= lambda x:x[1], reverse=True)[:10]
        string = "Topic "+str(i+1)+": "
        for t in sorted_terms:
            string = string + t[0] + ' '
            re = re + t[0] + ' '
        print(string)
    return re
```

TOPIC MODELING



POSITIVE



NEGATIVE

NEUTRAL COMMENT

neu_neg 3 stars comment with compound score \leq Q1

neu_pos 3 stars comment with compound score $>$ Q3

neu_nue 3 stars comment with compound between Q1 and Q3

neu_neg	2728		sentiment	Q1_compound	Q2_compound	Q3_compound
neu_neu	5457	0	negative	-0.7351	0.4019	0.9393
neu_pos	2721	1	neutral	0.7574	0.9459	0.9896
not_neu	89030	2	positive	0.9081	0.9689	0.9924

TOPIC MODELING



POSITIVE



NEGATIVE

TOPIC MODELING



A word cloud representing a topic related to a book store. The words are arranged in a vertical stack, with 'spent' being the largest word in the center. Other prominent words include 'bookstore', 'staff', and 'selection'. Smaller words like 'book', 'place', 'hours', 'time', 'store', 'new', 'powell', and 'portland' are also visible.

book place hours
spent
portland time store
bookstore
powell new
staff
selection

BOOK STORE



A word cloud representing a topic related to car rental. The words are arranged in a vertical stack, with 'free' being the largest word in the center. Other prominent words include 'booked', 'times', 'people', 'buy', and 'minutes'. Smaller words like 'scam', 'website', 'night', 'coming', 'charge', 'pick', 'price', 'available', 'shuttle', 'service', 'economy', 'experience', 'stars', 'read', 'instead', 'reviews', 'needed', 'paid', 'told', 'took', 'customer', 'orlando', 'late', 'waiting', 'airport', and 'online' are also visible.

took told paid online
day minutes
customer scam website orlando late
people buy airport
night coming charge pick
booked
reviews needed economy experience stars
read instead service shuttle price available
times

CAR RENTAL

CONCLUSION

- **Successfully extract why users give positive or negative comments**
 - **Extract the negative and positive comment from neutral comment**
 - **Understand why the neutral comment is skewed toward positive or negative**
 - **Offer accurate advice to the other business**
-

Q&A

